

Predicting Response with Decision Trees

Professor Song Yao
Olin Business School

Customer Analytics

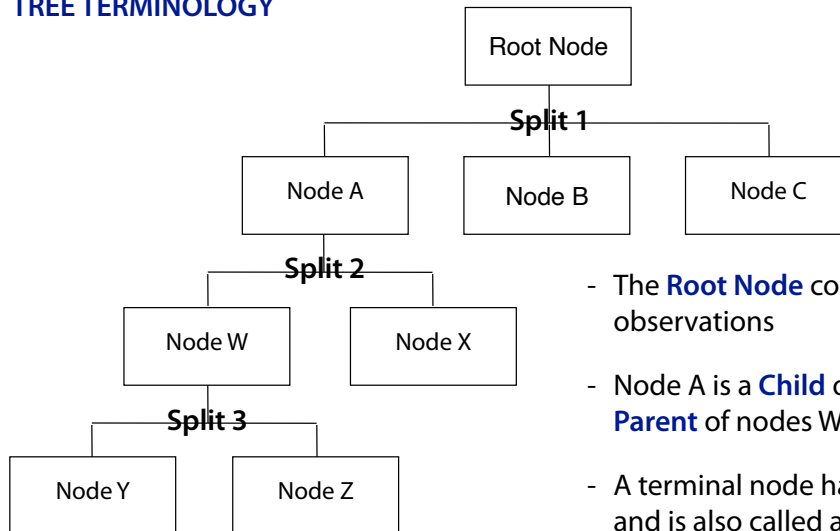
1

Basic Logic behind Decision Trees

2

How do we refer to the different parts of a tree?

TREE TERMINOLOGY



- The **Root Node** contains all observations
- Node A is a **Child** of the root and **Parent** of nodes W and X
- A terminal node has no children and is also called a **Leaf**
- There are **five leaves** B, C, X, Y, and Z (i.e., **five segments** of customers)

3

Fundamental concept of decision trees

CLASSIFICATION TREES

- The dependent variables is **categorical** (non-metric), e.g., buy/not buy

REGRESSION TREES

- The dependent variable is **numeric** (metric) , e.g. how much did a customer spend?

Basic Logic behind Decision Trees

- Each leaf (a segment of customers) in the tree is grown by splitting customers based on their independent variables (e.g., recency, gender, income, ...)
- The values of DV of customers of one leaf (e.g., buy or not buy, amount of spending) is different from those of customers in other leaves (as much as possible)

4

A Toy Example: First set of branches

5

We use a simple example to show how the CHAID (Chi-squared Automatic Interaction Detection) method works

- Predict response with customer demographics
- Dependent variable: Response -- "response"
- Independent variables:
 - Age -- "age"
 - Income -- "income"
 - Gender -- "female"

Response: DV

Response Distribution:		
	Count	Percentage (%)
response		
0	799	79.9
1	201	20.1

Income

Income Distribution:		
	Count	Percentage (%)
income		
1	208	20.8
2	487	48.7
3	305	30.5

Age

Age Distribution:		
	Count	Percentage (%)
age		
1	320	32.0
2	351	35.1
3	329	32.9

Female

Female Distribution:		
	Count	Percentage (%)
female		
0	517	51.7
1	483	48.3

6

CHAID EXAMPLE: ROOT NODE (TOP LEVEL)

Age variable:

- Cross tab every possible combination of two values of "age" with "response"
- Combine the two values that are the least significantly different from each other
- Stop if all remaining categories are significantly different in predicting response

```
# Compare age groups 1&2
age12 = chaid_demo[chaid_demo['age'].isin([1, 2])]
chi2_age12 = chi2_contingency(pd.crosstab(age12['response'],
                                          age12['age']), correction=False)

print(chi2_age12.pvalue)

# Compare age groups 1&3
age13 = chaid_demo[chaid_demo['age'].isin([1, 3])]
chi2_age13 = chi2_contingency(pd.crosstab(age13['response'],
                                          age13['age']), correction=False)

print(chi2_age13.pvalue)

# Compare age groups 2&3
age23 = chaid_demo[chaid_demo['age'].isin([2, 3])]
chi2_age23 = chi2_contingency(pd.crosstab(age23['response'],
                                          age23['age']), correction=False)

print(chi2_age23.pvalue)

7.84511863307372e-05
1.842310598122182e-06
0.3548097350546101
```

All possible combinations of different values of "age": 12, 13, 23

Age 1's response rates are significantly different from Age 2 and Age 3 (according to Chi2 test).

Treat age=2 and age=3 as equivalent in terms of predicting response
--> **combine them into one category**

7

CHAID EXAMPLE: ROOT NODE (TOP LEVEL)

Age variable: (take 2)

- Cross tab every possible combination of two values of "age" with "response"
 - Combine the two values that are the least significantly different from each other
 - ▶ Age 1 vs. Age 23
- Stop if all remaining categories are significantly different in predicting response

```
# Create new age grouping (1 vs 2&3 combined)
chaid_demo['ageNEW'] = np.where(chaid_demo['age'] == 1, 1, 23)
chi2_ageNEW = chi2_contingency(pd.crosstab(chaid_demo['response'],
                                          chaid_demo['ageNEW']), correction=False)

print(chi2_ageNEW.pvalue)

1.6626069931911872e-06
```

Age=1 and age=(2 or 3) are significant predictors of response
--> **Stop and go to next variable**

8

CHAID EXAMPLE: ROOT NODE (TOP LEVEL)

Income variable:

- Continue with "income" and "response": Possible combinations for income, 12, 13, 23
- Combine the two values that are the least significantly different from each other
- Stop if all remaining categories are significantly different in predicting response

```
# Compare income groups 1&2
income12 = chaid_demo[chaid_demo['income'].isin([1, 2])]
chi2_income12 = chi2_contingency(pd.crosstab(income12['response'],
                                             income12['income']), correction=False)
print(chi2_income12.pvalue)

# Compare income groups 1&3
income13 = chaid_demo[chaid_demo['income'].isin([1, 3])]
chi2_income13 = chi2_contingency(pd.crosstab(income13['response'],
                                             income13['income']), correction=False)
print(chi2_income13.pvalue)

# Compare income groups 2&3
income23 = chaid_demo[chaid_demo['income'].isin([2, 3])]
chi2_income23 = chi2_contingency(pd.crosstab(income23['response'],
                                             income23['income']), correction=False)
print(chi2_income23.pvalue)
```

0.10568540009415126
0.04793679612791425
1.4713066443738882e-05

Treat income=1 and income=2 as
equivalent in terms of predicting
response (Chi2 test insignificant
--> combine them into
1 category

9

CHAID EXAMPLE: ROOT NODE (TOP LEVEL)

Income variable: (take 2)

- Cross tab every possible combination of two values of "income" with "response"
 - Combine the two values that are the least significantly different from each other
 - ▶ Income 12 and Income 3
- Stop if all remaining categories are significantly different in predicting response

```
# Create new income grouping (1&2 combined vs 3)
chaid_demo['incomeNEW'] = np.where(chaid_demo['income'].isin([1, 2]), 12, 3)
chi2_incomeNEW = chi2_contingency(pd.crosstab(chaid_demo['response'],
                                              chaid_demo['incomeNEW']), correction=False)
print(chi2_incomeNEW.pvalue)
```

4.884343487542051e-05

income=(1 or 2) and
income=3 are significant
predictors of response
--> Stop and go to
next variable

10

CHAID EXAMPLE: ROOT NODE (TOP LEVEL)

Female variable:

- Cross tab every possible combination of two values of "female" with "response"
- Combine the two values that are the least significantly different from each other
- Stop if all remaining categories are significantly different in predicting response

```
chi2_female = chi2_contingency(pd.crosstab(chaid_demo['response'],
                                           chaid_demo['female']), correction=False)
print(chi2_female.pvalue)

### based on the p values, female is the most significant variable as it has the lowest p value
### we will use female to split the data this round
```

2.0152561677568327e-07



female=0 and female=1
are significant predictors
of response
-->
**Stop and proceed to
selecting the variable to
first split the sample**

11

CHAID EXAMPLE: ROOT NODE (TOP LEVEL)

Select variable for first sample partition:

- Cross tab every final variable (after combining categories) with "response"
- Select the variable with the smallest p-value

```
chi2_female = chi2_contingency(pd.crosstab(chaid_demo['response'],
                                           chaid_demo['female']), correction=False)
print(chi2_female.pvalue)

### based on the p values, female is the most significant variable as it has the lowest p value
### we will use female to split the data this round
```

2.0152561677568327e-07

```
# Create new age grouping (1 vs 2&3 combined)
chaid_demo['ageNEW'] = np.where(chaid_demo['age'] == 1, 1, 23)
chi2_ageNEW = chi2_contingency(pd.crosstab(chaid_demo['response'],
                                           chaid_demo['ageNEW']), correction=False)
print(chi2_ageNEW.pvalue)
```

1.6626069931911872e-06

```
# Create new income grouping (1&2 combined vs 3)
chaid_demo['incomeNEW'] = np.where(chaid_demo['income'].isin([1, 2]), 12, 3)
chi2_incomeNEW = chi2_contingency(pd.crosstab(chaid_demo['response'],
                                              chaid_demo['incomeNEW']), correction=False)
print(chi2_incomeNEW.pvalue)
```

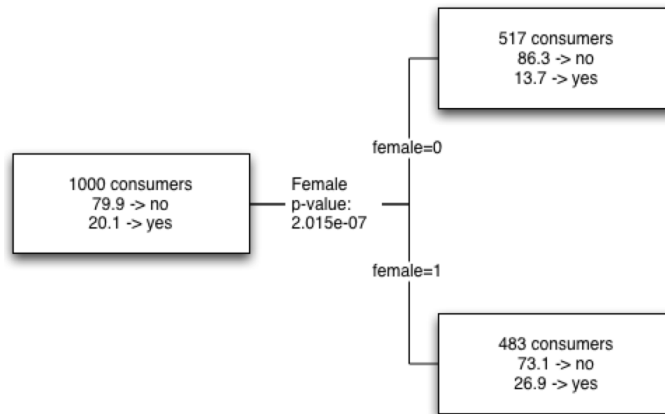
4.884343487542051e-05

Female split has the smallest p-value
Pick gender for sample partition at root node

12

We have determined that gender is the most important predictor

CHAID EXAMPLE: FIRST DATA PARTITION



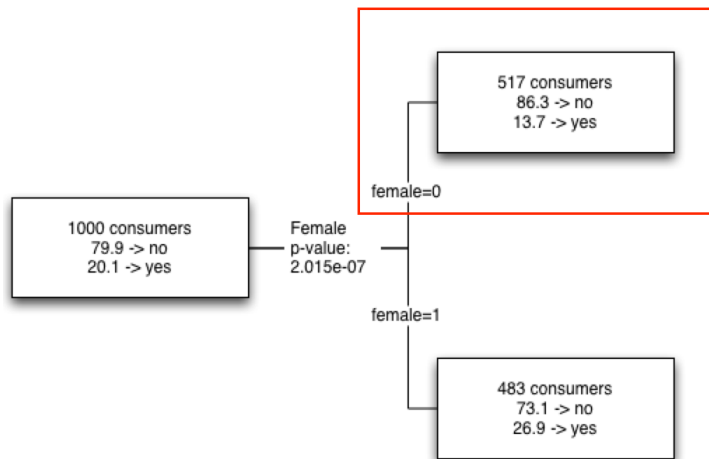
13

**A Toy Example Continued:
Growing the Tree**

14

We proceed with the female=0 (i.e., male) branch

CHAID EXAMPLE: FIRST DATA PARTITION



15

CHAID EXAMPLE: FIRST CHILD NODE (where FEMALE=0)

Age variable:

- Cross tab every possible combination of two values of "age" with "response"
- Combine the two values that are the least significantly different from each other
- Stop if all remaining categories are significantly different in predicting response

```
# Compare age groups 1&2 for males
male_age12 = chaid_demo[(chaid_demo['age'].isin([1, 2])) &
                        (chaid_demo['female'] == 0)].copy()
chi2_male_age12 = chi2_contingency(pd.crosstab(male_age12['response'],
                                                male_age12['age']), correction=False)
print(chi2_male_age12.pvalue)

# Compare age groups 1&3 for males
male_age13 = chaid_demo[(chaid_demo['age'].isin([1, 3])) &
                        (chaid_demo['female'] == 0)].copy()
chi2_male_age13 = chi2_contingency(pd.crosstab(male_age13['response'],
                                                male_age13['age']), correction=False)
print(chi2_male_age13.pvalue)

# Compare age groups 2&3 for males
male_age23 = chaid_demo[(chaid_demo['age'].isin([2, 3])) &
                        (chaid_demo['female'] == 0)].copy()
chi2_male_age23 = chi2_contingency(pd.crosstab(male_age23['response'],
                                                male_age23['age']), correction=False)
print(chi2_male_age23.pvalue)
```

0.36734481098374727

0.2117626995367813

0.7097908825361634

Treat age=2 and age=3 as equivalent in terms of predicting response
--> combine them into 1 category

16

CHAID EXAMPLE: FIRST CHILD NODE (FEMALE=0)

Age variable: (take 2)

- Cross tab every possible combination of two values of "age" with "response"
- Combine the two values that are the least significantly different from each other
- Stop if all remaining categories are significantly different in predicting response

```
# Create new age grouping (1 vs 2&3 combined) for males
male_data = chaid_demo[chaid_demo['female'] == 0].copy()
male_data['ageNEW'] = np.where(male_data['age'] == 1, 1, 23)
chi2_male_ageNEW = chi2_contingency(pd.crosstab(male_data['response'],
                                                male_data['ageNEW']), correction=False)
print(chi2_male_ageNEW.pvalue)
```

0.22786305389769565

Age=1 and age=(2 or 3)
are **not** significant
predictors of response
--> **Combine all age
categories, i.e. ignore
age as predictor for
males!**

**Stop and go to next
variable**

17

CHAID EXAMPLE: FIRST CHILD NODE (FEMALE=0)

Income variable:

- Cross tab every possible combination of two values of "income" with "response"
- Combine the two values that are the least significantly different from each other
- Stop if all remaining categories are significantly different in predicting response

```
# Compare income groups 1&2 for males
male_income12 = chaid_demo[(chaid_demo['income'].isin([1, 2])) &
                             (chaid_demo['female'] == 0)].copy()
chi2_male_income12 = chi2_contingency(pd.crosstab(male_income12['response'],
                                                  male_income12['income']), correction=False)
print(chi2_male_income12.pvalue)
```

```
# Compare income groups 1&3 for males
male_income13 = chaid_demo[(chaid_demo['income'].isin([1, 3])) &
                             (chaid_demo['female'] == 0)].copy()
chi2_male_income13 = chi2_contingency(pd.crosstab(male_income13['response'],
                                                  male_income13['income']), correction=False)
print(chi2_male_income13.pvalue)
```

```
# Compare income groups 2&3 for males
male_income23 = chaid_demo[(chaid_demo['income'].isin([2, 3])) &
                             (chaid_demo['female'] == 0)].copy()
chi2_male_income23 = chi2_contingency(pd.crosstab(male_income23['response'],
                                                  male_income23['income']), correction=False)
print(chi2_male_income23.pvalue)
```

0.933515557879672

0.042925591804953415

0.005081143114016782

Treat income=1 and income=2 as equivalent
in terms of predicting response
--> **combine them into 1 category**

18

CHAID EXAMPLE: FIRST CHILD NODE (FEMALE=0)

Income variable: (take 2)

- Cross tab every possible combination of two values of "income" with "response"
- Combine the two values that are the least significantly different from each other
- Stop if all remaining categories are significantly different in predicting response

```
# Create new income grouping (1&2 combined vs 3) for males
male_data = chaid_demo[chaid_demo['female'] == 0].copy()
male_data['incomeNEW'] = np.where(male_data['income'].isin([1, 2]), 12, 3)
chi2_male_incomeNEW = chi2_contingency(pd.crosstab(male_data['response'],
                                                    male_data['incomeNEW']), correction=False)

print(chi2_male_incomeNEW.pvalue)
```

0.0026573404504472826

income=(1 or 2) and
income=3 are significant
predictors of response
--> **Stop and proceed to
choose which variable to
split the sample next**

19

CHAID EXAMPLE: FIRST CHILD NODE (FEMALE=0)

Select variable for second sample partition (in the female=0 branch):

- Cross tab every final variable (after combining categories) with "response"
- Select the variable with the smallest p-value

```
# Create new income grouping (1&2 combined vs 3) for males
male_data = chaid_demo[chaid_demo['female'] == 0].copy()
male_data['incomeNEW'] = np.where(male_data['income'].isin([1, 2]), 12, 3)
chi2_male_incomeNEW = chi2_contingency(pd.crosstab(male_data['response'],
                                                    male_data['incomeNEW']), correction=False)

print(chi2_male_incomeNEW.pvalue)
```

0.0026573404504472826

```
# Create new age grouping (1 vs 2&3 combined) for males
male_data = chaid_demo[chaid_demo['female'] == 0].copy()
male_data['ageNEW'] = np.where(male_data['age'] == 1, 1, 23)
chi2_male_ageNEW = chi2_contingency(pd.crosstab(male_data['response'],
                                                male_data['ageNEW']), correction=False)

print(chi2_male_ageNEW.pvalue)
```

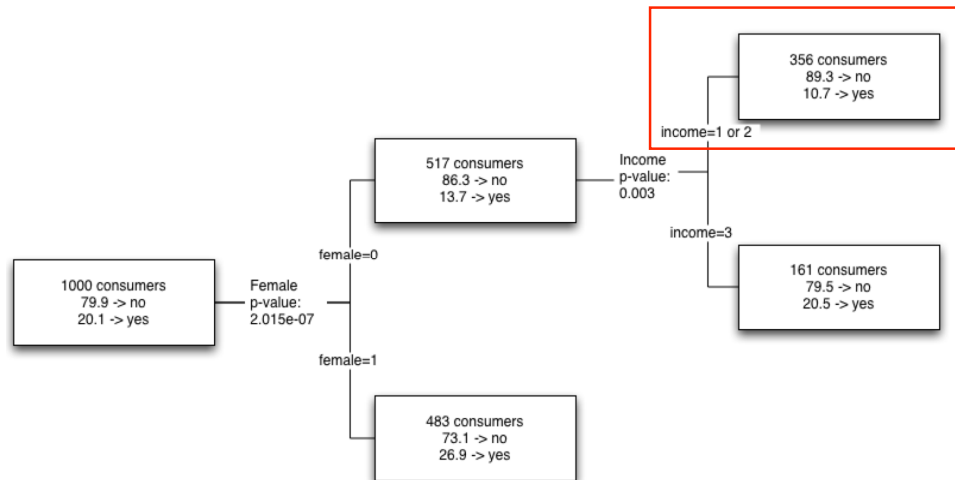
0.22786305389769565

Smallest p-value
-->
**pick income for sample
partition in the female=0
branch**

20

We have determined that income is the most important predictor for men

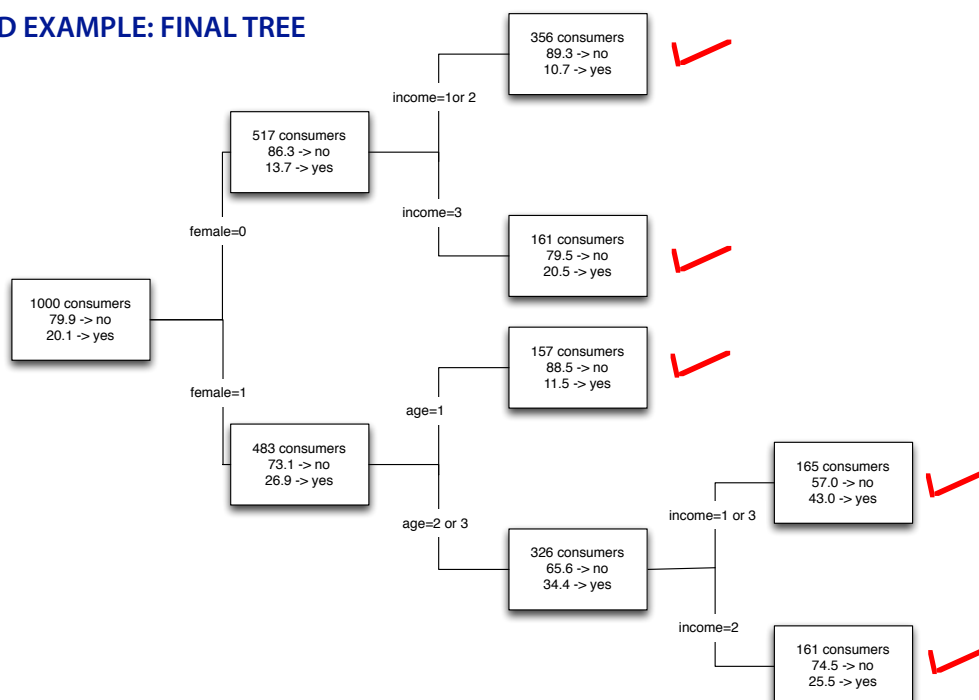
CHAID EXAMPLE: THIRD DATA PARTITION



21

One more split completes the tree

CHAID EXAMPLE: FINAL TREE



22

How do we use Decision Trees to make decision

23

Decision trees are built through recursive partitioning

IDEA OF DECISION TREES ALGORITHMS

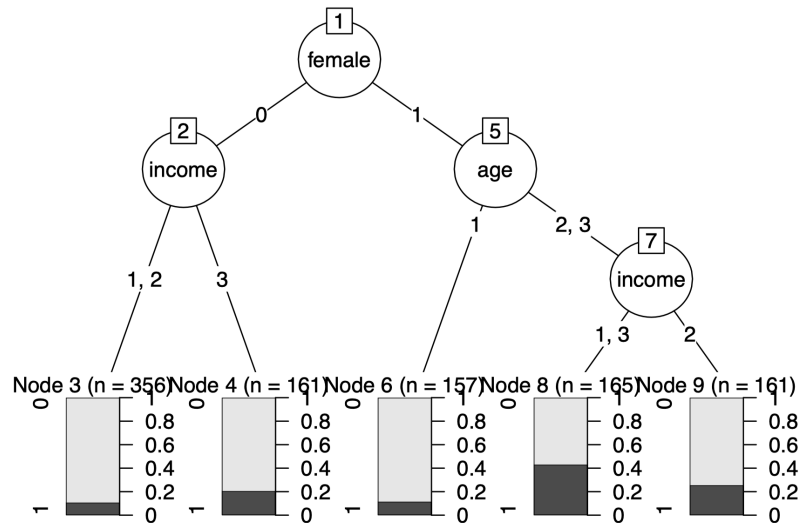
- Recursive partitioning = an iterative process of splitting the data into partitions and then splitting each partition into a sub-partitions
- Partition = MECE (**M**utually **E**xclusive and **C**ollectively **E**xhaustive)
- Start with Root node (includes all observations in the training sample)
- Each branch finds a single variable to split the data in 2 or more groups
 - Algorithm tries to break up the data, often using every possible split on every variable (brute force --> may require huge computing power)
 - For binary tree, suppose ages run 18-94 and consider splitting on $\leq 18 / > 18$ versus $\leq 19 / > 19$ versus $\leq 20 / > 20$ versus...versus $\leq 93 / > 93$
 - The best split is the one which best separates groups as measured by the dependent variable
- We focused on CHAID as an example. There are many different implementations of this basic algorithm (e.g., CHAID, CART, C4.5, etc.).

24

R and Python both have functions for CHAID, but they are somewhat cumbersome (particularly Python)

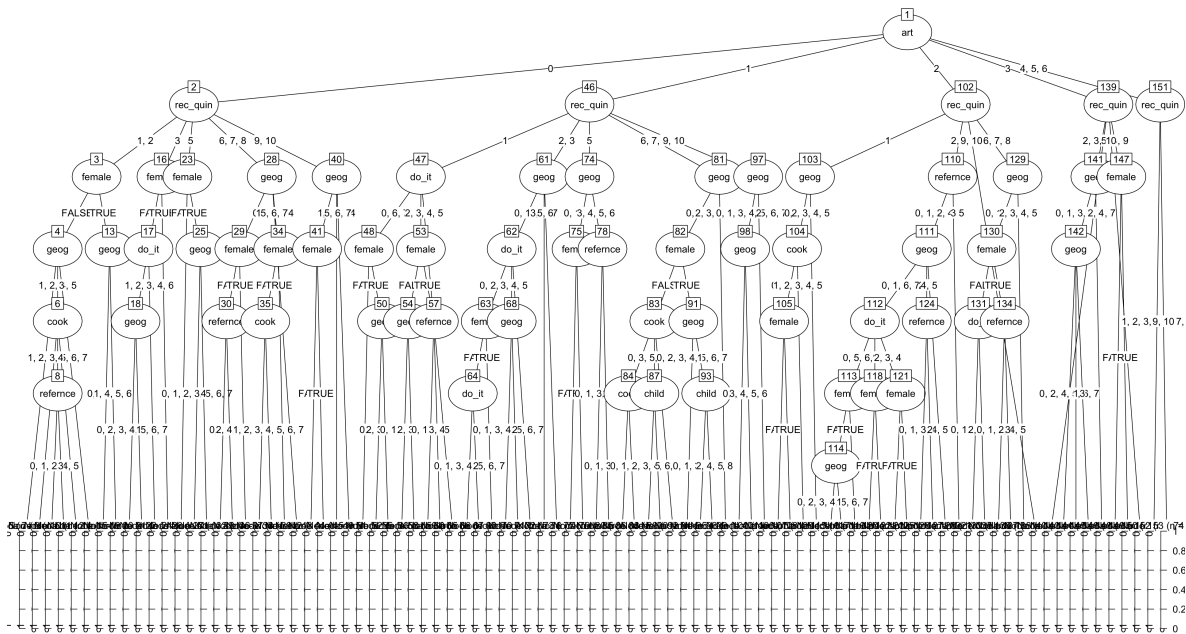
CHAID EXAMPLE: FINAL TREE

```
chaid_demo_model <- chaid(response ~ age + female + income, data = chaid_demo)
plot(chaid_demo_model)
```



25

CHAID TREE FOR BOOKBINDERS BUYERS



26

The break-even response rate tells us to which cells to extend the offer

BREAK EVEN RESPONSE RATE

- Cost of mailing an offer = \$0.50
- Selling price (includes shipping) = \$18
- Wholesale price paid by Bookbinders = \$9
- Shipping costs = \$3
- Break-even = Cost to mail/net revenue per sale =
 $.5/(18-9-3) = 8.3\%$
- **Depending which leaves (segments) have response rates greater than 8.3%, we decide how to target.**

27

We can use decision tree based analysis for a range of purposes

USES OF DECISION TREES

- **Prediction/classification:**
Like logistic regression, predict values of a target variable
- **Segmentation:**
Identify relatively homogeneous groups
- **Interaction identification:**
Identify relationships that pertain only to specific subgroups, for example for use in a logistic regression model

28

Alternative Decision Tree Method to CHAID

CART, Classification and Regression Tree

- **Classification tree:**
Instead of using Chi2, Gini-index is used for splitting branches
 - Gini-index: Another metric for the difference of DV across nodes
- **Regression tree:**
ANOVA is one example metric that can be used for splitting branches
- **K-fold cross validation automatically**

29

Decision trees are simple to understand and implement

ADVANTAGES OF DECISION TREES

- Actionable -- generates a set of simple rules which can be used to classify/predict new cases (e.g. if age<25 and purchase at least 1 art book last year, will purchase "The Art History of Florence")
 - Easy to encode rules in decision systems
 - Fast to classify and predict new cases
- Provides a clear indication which variables are most important for prediction or classification

30

Random Forest

31

Decision trees also have some serious disadvantages

DISADVANTAGES OF DECISION TREES

- 'Lose' information compared to regression models b/c of categorization of continuous variables:
 - In regression model, the prediction will be different for each value of a continuous predictor, e.g. 1014 vs. 1016 vs. 2006 vs. 2100
 - In decision tree, 1014 and 1016 are likely to fall into the same node and thus have the same prediction.
- Trees can be too large to properly interpret
- Can be error-prone if the number of observations per class gets small
- Can fit well training sample but badly in test sample -- overfitting problem

32

“Ensembles” of trees produce better predictions and are less prone to overfitting

ENSEMBLE APPROACHES

- Random Forests
- Boosted Decision Trees

CORE IDEA

- Each tree *may be* a weak predictor
- Create many decision trees
 - using subsets of the original data and subsets of the independent variables (Random Forests)
 - weighting original data differently (Boosted Decision Trees)
- Go with the (weighted) average of the prediction of all the decision trees

33

“Random Forest” is the most popular type of decision tree ensemble

RANDOM FOREST IDEA

- Decision trees algorithm that injects randomness into fitting (invented by Leo Breiman)
- This randomness reduces overfitting
- Key idea is to create many decision trees (~500), each of them based on
 - randomly chosen subsample of the data
 - randomly chosen subset of the predictor variables at each node that is considered
- Very accurate predictor, can handle huge number of input variables

34

The random forest algorithm builds hundreds of trees

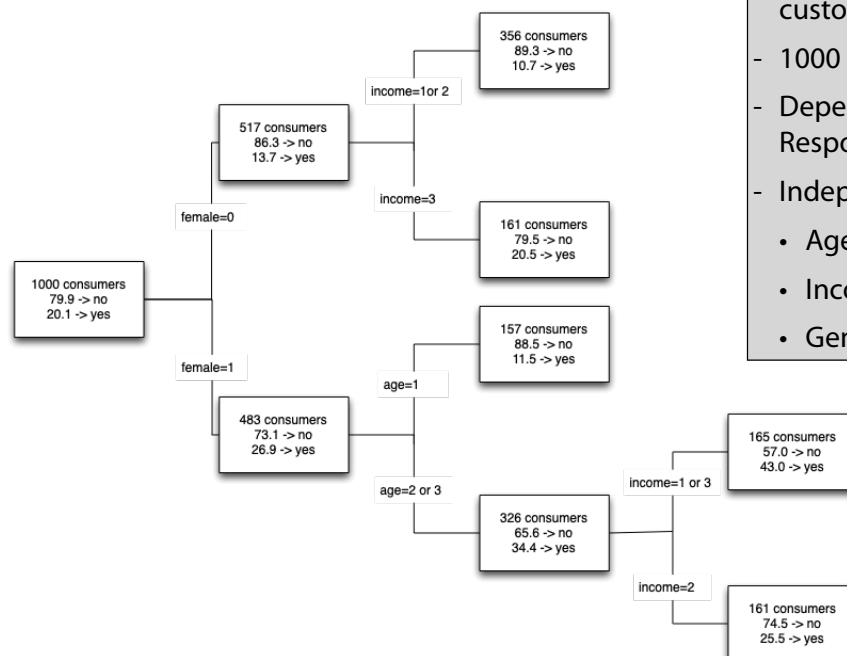
RANDOM FOREST ALGORITHM

- For each tree:
 - Create a “Bootstrap Sample”
 - Let’s say the original sample has N observations
 - Randomly draw observations from the original data (with replacement) and form a new data of N observations
 - Some observations may appear multiple times; some may never appear
 - Sample the variables when splitting
 - When building a branch, randomly select a subset of variables for the split instead of considering all available variables
 - The number of variables in the subset is typically the square root of the number of independent variables (e.g. 50 variables, randomly pick 7)
 - At next node, randomly select another subset of variables
- “Ensemble Scoring”
 - Combine the results from all the trees by averaging the prediction of each tree

35

Recall the CHAID example

CHAID EXAMPLE: FINAL TREE



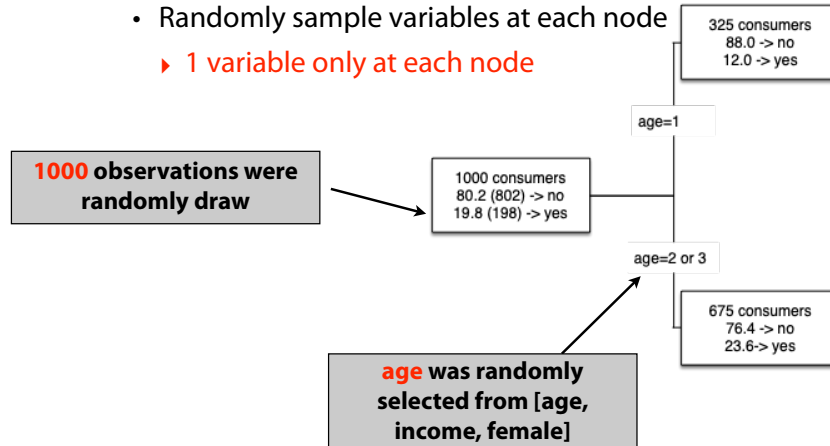
- Predict response with customer demographics
- 1000 consumers
- Dependent variable: Response -- “response”
- Independent variables:
 - Age -- “age”
 - Income -- “income”
 - Gender -- “female”

36

Lets predict using Random Forest

RANDOM FOREST EXAMPLE

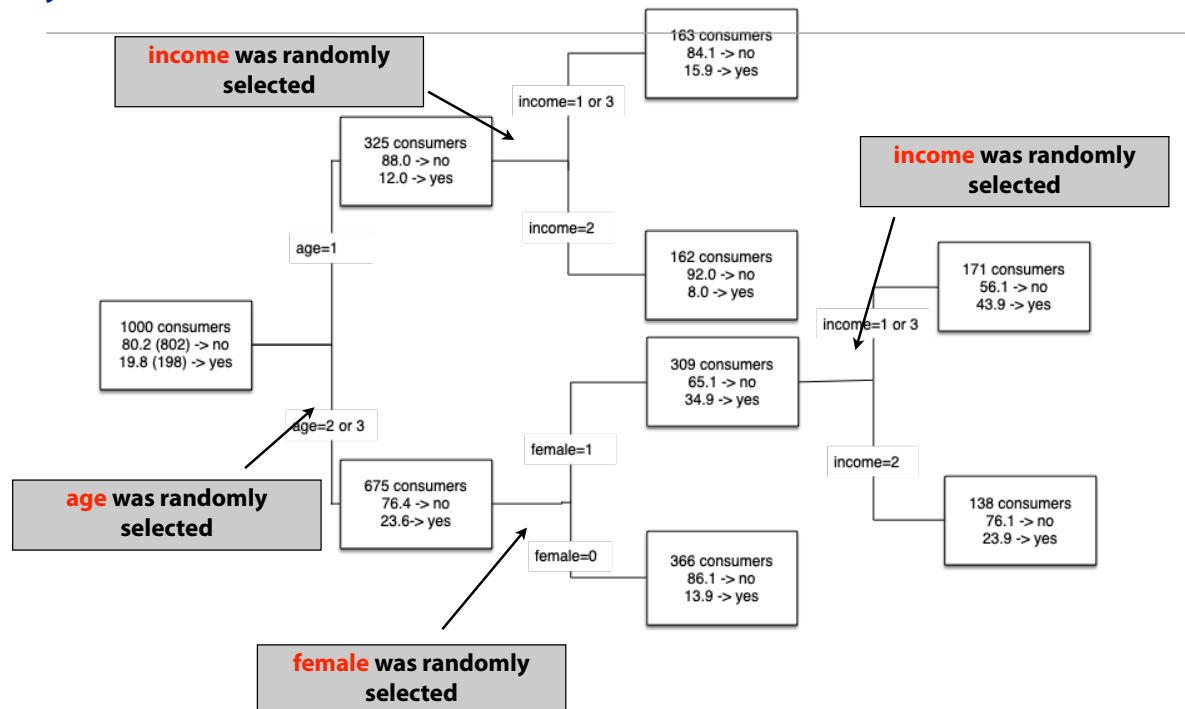
- For each tree:
 - "Bootstrap" observations
 - ▶ Draw 1000 observations from the original dataset (with replacement)
 - Randomly sample variables at each node
 - ▶ 1 variable only at each node



- Predict response with customer demographics
- 1000 consumers
- Dependent variable: Response -- "response"
- Independent variables:
 - Age -- "age"
 - Income -- "income"
 - Gender -- "female"

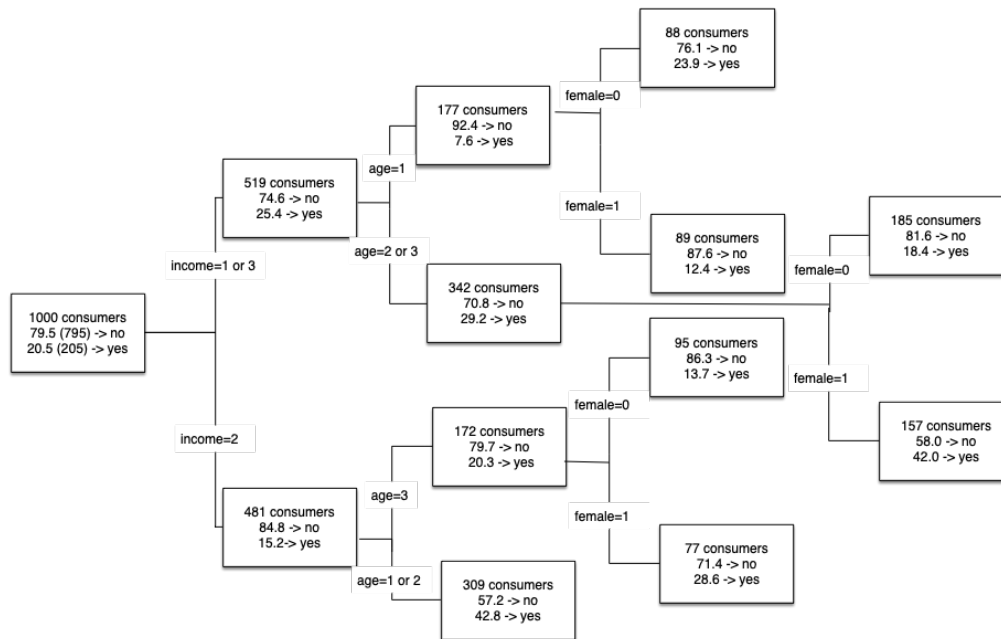
37

Repeating the random selection of variables at each node yields the full **FIRST** tree



38

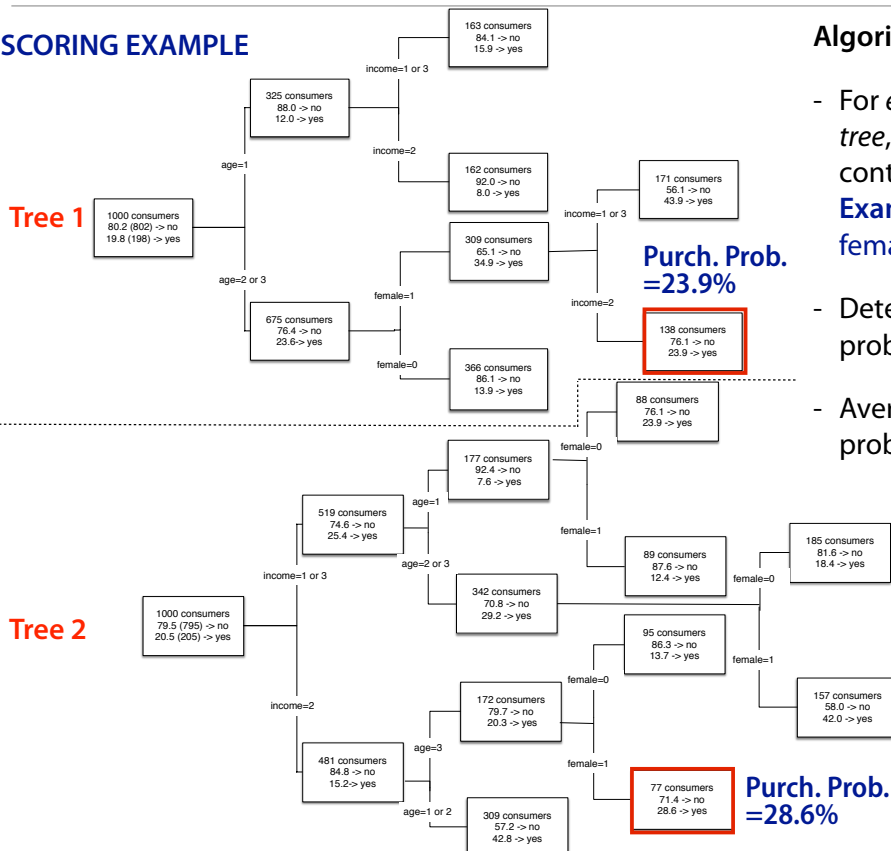
Resampling observations and resampling variables at each node yields the full **SECOND** tree



39

We can now “score the ensemble” of trees

SCORING EXAMPLE



Algorithm

- For each customer and each tree, determine the leaf that contains that customer
Example: A customer female=1, age=3, income=2
- Determine the purchase probability for each leaf
- Average the purchase probabilities over the trees
- **Predicted Purchase Probability of This Customer**

$$= (23.9\% + 28.6\%) / 2 = 26.25\%$$

40

In practice we would construct and ensemble score hundreds of trees

- Instead of just having two trees, build, say, 500 of them.
 - Given a customer's age, income, and gender, the customer will be grouped into a particular leaf at each respective tree.
 - Each leaf has its different response rate prediction for that customer (e.g., 23.9% vs. 28.6%)
 - Each customer has 500 response rates from the 500 leaves
 - The average of these 500 response rates is the predicted response probability of the customer

41

Random Forests: Advantages and Disadvantages

ADVANTAGES

- Very accurate predictor
 - Random forests algorithm is less prone to overfitting (not completely)
 - ▶ Because of the randomness, observations and variables that may cause overfitting are always dropped at some point
 - With a single decision tree, when alternative variables have similar prediction power, it may be difficult to decide how to build a node
 - ▶ Random Forest allows such variables all contribute to the prediction because we randomly select variables at each node (each one of them has the chance to be picked)
- Can handle huge number of input variables
 - Only use a subset of variables at each node

DISADVANTAGES

- Harder to interpret the results than single decision tree
- Time consuming
 - Hundreds of trees

42

Many models—need to tune and compare carefully

Firewall Example: Model Comparison on the Test Sample ([*Code on Canvas for your ref.*](#))

